

Random Forest and Decision Trees- Loan data

March 31, 2021

Random Forest and Decision Trees on Loan Data

Sohini Mukherjee

31.03.2021

Importing Libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Getting the Data

```
[3]: loans = pd.read_csv('E:/2.PYTHON-ML-BOOTCAMP/resources/
↳15-Decision-Trees-and-Random-Forests/loan_data.csv')
```

```
[4]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose                9578 non-null   object
2   int.rate               9578 non-null   float64
3   installment            9578 non-null   float64
4   log.annual.inc         9578 non-null   float64
5   dti                    9578 non-null   float64
6   fico                   9578 non-null   int64
7   days.with.cr.line      9578 non-null   float64
8   revol.bal              9578 non-null   int64
9   revol.util             9578 non-null   float64
10  inq.last.6mths         9578 non-null   int64
11  delinq.2yrs            9578 non-null   int64
12  pub.rec                9578 non-null   int64
13  not.fully.paid         9578 non-null   int64
```

```
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

```
[5]: loans.describe()
```

```
[5]:      credit.policy    int.rate  installment  log.annual.inc      dti  \
count      9578.000000  9578.000000  9578.000000      9578.000000  9578.000000
mean         0.804970    0.122640   319.089413      10.932117    12.606679
std          0.396245    0.026847   207.071301       0.614813     6.883970
min          0.000000    0.060000    15.670000       7.547502     0.000000
25%          1.000000    0.103900   163.770000      10.558414     7.212500
50%          1.000000    0.122100   268.950000      10.928884    12.665000
75%          1.000000    0.140700   432.762500      11.291293    17.950000
max          1.000000    0.216400   940.140000      14.528354    29.960000

      fico  days.with.cr.line    revol.bal    revol.util  \
count  9578.000000      9578.000000  9.578000e+03  9578.000000
mean   710.846314    4560.767197   1.691396e+04   46.799236
std    37.970537    2496.930377   3.375619e+04   29.014417
min    612.000000    178.958333   0.000000e+00    0.000000
25%    682.000000    2820.000000   3.187000e+03   22.600000
50%    707.000000    4139.958333   8.596000e+03   46.300000
75%    737.000000    5730.000000   1.824950e+04   70.900000
max    827.000000   17639.958330   1.207359e+06  119.000000

      inq.last.6mths  delinq.2yrs    pub.rec  not.fully.paid
count      9578.000000  9578.000000  9578.000000      9578.000000
mean         1.577469    0.163708    0.062122      0.160054
std          2.200245    0.546215    0.262126      0.366676
min          0.000000    0.000000    0.000000      0.000000
25%          0.000000    0.000000    0.000000      0.000000
50%          1.000000    0.000000    0.000000      0.000000
75%          2.000000    0.000000    0.000000      0.000000
max          33.000000   13.000000    5.000000      1.000000
```

```
[6]: loans.head()
```

```
[6]:      credit.policy      purpose  int.rate  installment  log.annual.inc  \
0           1  debt_consolidation    0.1189         829.10     11.350407
1           1      credit_card    0.1071         228.22     11.082143
2           1  debt_consolidation    0.1357         366.86     10.373491
3           1  debt_consolidation    0.1008         162.34     11.350407
4           1      credit_card    0.1426         102.92     11.299732

      dti  fico  days.with.cr.line    revol.bal    revol.util  inq.last.6mths  \
0  19.48  737    5639.958333    28854        52.1           0
1  14.29  707    2760.000000    33623        76.7           0
2  11.63  682   4710.000000     3511        25.6           1
```

3	8.10	712	2699.958333	33667	73.2	1
4	14.97	667	4066.000000	4740	39.5	0


	delinq.2yrs	pub.rec	not.fully.paid
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	1	0	0

Exploratory Data Analysis

Creating a histogram of two FICO distributions on top of each other, one for each credit.policy outcome.

```
[7]: plt.figure(figsize=(10,5))

loans[loans['credit.policy']==1]['fico'].hist(alpha=0.5, bins=30, color='blue',
→label='Credit.Policy=1', ec='black')
loans[loans['credit.policy']==0]['fico'].hist(alpha=0.5, bins=30, color='red',
→label='Credit.Policy=0', ec='black')
plt.legend()
plt.show()
```




Creating a histogram of two FICO distributions on top of each other, selected by the not.fully.paid column.

```
[8]: plt.figure(figsize=(10,5))
```

```

loans[loans['not.fully.paid']==1]['fico'].hist(alpha=0.5, bins=30,
→color='blue', label='Credit.Policy=1', ec='black')
loans[loans['not.fully.paid']==0]['fico'].hist(alpha=0.5, bins=30, color='red',
→label='Credit.Policy=0', ec='black')
plt.legend()
plt.show()

```



Creating a countplot using seaborn showing the counts of loans by purpose, with the color hue defined by not.fully.paid.

```

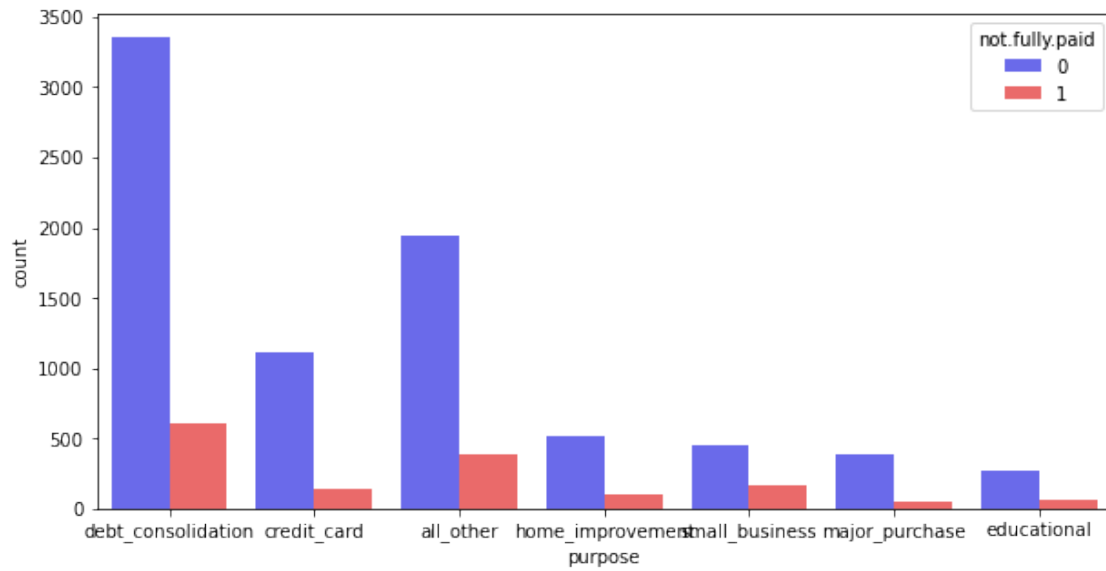
[9]: plt.figure(figsize=(10,5))
     sns.countplot(x='purpose', data=loans, hue='not.fully.paid', palette='seismic')
     plt.tight_layout

```

```

[9]: <function matplotlib.pyplot.tight_layout(*, pad=1.08, h_pad=None, w_pad=None,
     rect=None)>


```



Creating the trend between FICO score and interest rate.

```
[10]: sns.jointplot(x='fico', y='int.rate', data=loans, color='purple')
```


```
[10]: <seaborn.axisgrid.JointGrid at 0x1fab3b365b0>
```



Creating lmplots to see if the trend differed between not.fully.paid and credit.policy

```
[11]: sns.lmplot(x='fico', y='int.rate', data=loans, col='not.fully.paid',
               hue='credit.policy', palette='seismic')
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x1fab3d53ee0>
```



Since we have ‘purpose’ as a categorical variable we need to make dummies for the column for ML to work

```
[12]: cat_feats = ['purpose']
final_data= pd.get_dummies(loans, columns=cat_feats, drop_first= True)
final_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 9578 entries, 0 to 9577
```

```
Data columns (total 19 columns):
```

#	Column	Non-Null Count	Dtype
0	credit.policy	9578 non-null	int64
1	int.rate	9578 non-null	float64
2	installment	9578 non-null	float64
3	log.annual.inc	9578 non-null	float64
4	dti	9578 non-null	float64
5	fico	9578 non-null	int64
6	days.with.cr.line	9578 non-null	float64
7	revol.bal	9578 non-null	int64
8	revol.util	9578 non-null	float64
9	inq.last.6mths	9578 non-null	int64
10	delinq.2yrs	9578 non-null	int64
11	pub.rec	9578 non-null	int64
12	not.fully.paid	9578 non-null	int64
13	purpose_credit_card	9578 non-null	uint8
14	purpose_debt_consolidation	9578 non-null	uint8
15	purpose_educational	9578 non-null	uint8
16	purpose_home_improvement	9578 non-null	uint8
17	purpose_major_purchase	9578 non-null	uint8
18	purpose_small_business	9578 non-null	uint8

dtypes: float64(6), int64(7), uint8(6)
memory usage: 1.0 MB

Train Test Split

```
[13]: from sklearn.model_selection import train_test_split

X= final_data.drop('not.fully.paid', axis=1)
y= final_data['not.fully.paid']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=101)
```

Training a Decision Tree Model

```
[14]: from sklearn.tree import DecisionTreeClassifier

dtree = DecisionTreeClassifier()

dtree.fit(X_train, y_train)
```

```
[14]: DecisionTreeClassifier()
```

Predictions and Evaluation of Decision Tree

```
[15]: predictions = dtree.predict(X_test)

from sklearn.metrics import classification_report, confusion_matrix

print(classification_report(y_test, predictions))
print('\n')
confusion_matrix(y_test, predictions)
```

	precision	recall	f1-score	support
0	0.85	0.82	0.84	2431
1	0.19	0.23	0.21	443
accuracy			0.73	2874
macro avg	0.52	0.53	0.52	2874
weighted avg	0.75	0.73	0.74	2874

```
[15]: array([[1996, 435],
[ 340, 103]], dtype=int64)
```

Training the Random Forest model


```
[17]: from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier(n_estimators=200)

rfc.fit(X_train, y_train)
```

```
[17]: RandomForestClassifier(n_estimators=200)
```

Predictions and Evaluation

```
[18]: pred = rfc.predict(X_test)

print(classification_report(y_test, pred))
print('\n')
confusion_matrix(y_test, pred)
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	2431
1	0.53	0.02	0.04	443
accuracy			0.85	2874
macro avg	0.69	0.51	0.48	2874
weighted avg	0.80	0.85	0.78	2874

```
[18]: array([[2422,    9],
           [ 433,   10]], dtype=int64)
```

Random Forest gives better accuracy than Decision Trees